

What is Philosophy?¹

1. In this essay, I want to address the question: "What is philosophy?"

We might to begin with acknowledge a distinction between things that have *natures* and things that have *histories*. Physical things such as electrons and aromatic compounds would be paradigmatic of the first class, while cultural formations such as English Romantic poetry and Ponzi schemes would be paradigmatic of the second. Applied to the case at hand, this distinction would surely place philosophy on the side of things that have histories. But now we might ask: Does philosophy differ in this respect from physics, chemistry, or biology? Physical, chemical, and biological *things* have natures rather than histories, but what about the disciplines that define and study them? Should physics itself be thought of as something that has a nature, or as something that has a history? Concluding the latter is giving a certain kind of pride of place to the historical. For it is in effect treating the *distinction* between things that have natures and things that have histories, between things studied by the *Naturwissenschaften* and things studied by the *Geisteswissenschaften*, as itself a cultural formation: the sort of thing that itself has a history rather than a nature. And from here it is a short step (though not, to be sure, an obligatory one) to the thought that natures themselves are the sort of thing that have a history; certainly the *concepts* electron and aromatic compound are that sort of thing. At this point the door is opened to a thorough-going historicism. It is often thought that this is the point to which Hegel—one of my particular heroes—brought us. I think that thought is correct, as far as it goes, but that we go very wrong if we think that that is where Hegel left us.

To say that philosophy is, at least to begin with, to be understood as the sort of thing that has a history rather than a nature is to foreground the way in which what deserves to be counted as distinctively philosophical activity answers to what has actually been done by those we recognize as precedential, tradition-transforming philosophers. One of Hegel's deepest and most important insights, I think, is indeed that the determinate contentfulness of any universal—in this case, the concept of philosophy—can only be understood in terms of the process by which it incorporates the contingencies of the particulars to which it has actually been applied. But he goes on from there to insist that it is in each case the responsibility of those of us who are heirs to such a conceptual tradition to see to it that is a *rational* tradition—that the distinction it embodies and enforces between correct and incorrect applications of a concept can be *justified*, that applying it in one case and withholding application in another is something for which *reasons* can be given. It is only insofar as we can do that that we are entitled to understand what we are doing as applying *concepts*. We fulfill that obligation by rationally reconstructing the tradition, finding a coherent, cumulative trajectory through it that reveals it as expressively progressive—as

¹This chapter was originally published under the title "Reason, Expression, and the Philosophic Enterprise," in *What Is Philosophy?*, C.P. Ragland and Sarah Heidt (eds.), Yale University Press, 2001, pp. 74-95.

the gradual unfolding into greater explicitness of commitments that can be seen retrospectively as always already having been implicit in it. That is, it is our job to rewrite the history so as to discover in it the revelation of what then retrospectively appears as an antecedent nature. Hegel balances the insight that even natures have histories by seeing rationality itself as imposing the obligation to construe histories as revelatory of natures.

The aim is to pick out a sequence of precedential instances or applications of a concept that amount to the delineation of a content for the concept, much as a judge at common law is obliged to do. *Making* the tradition rational, is not independent of the labor of concretely *taking* it to be so. It is a criterion of adequacy of each such Whiggish rewriting of our disciplinary history that it create and display continuity and progress by its systematic inclusions and exclusions. The discontinuities that correspond to shifts of topic, the forgetting of lessons, and the degeneration of research programs are invisible from within each such telling; but those differences live on in the spaces between the tellings. Each generation redefines its subject by offering a new retrospective reading of its characteristic concerns and hard-won lessons.² But also, at any one time there will be diverse interpretations, complete with rival canons, competing designations of heroes, and accounts of their heroic feats. Making canons and baking traditions out of the rich ingredients bequeathed us by our discursive predecessors is a game that all can play.

In this chapter, I am going to sketch one such perspective on what philosophers do—discern a nature as revealed by the history.

Ours is a broadly cognitive enterprise—I say ‘*broadly* cognitive’ to indicate that I mean that philosophers aim at a kind of *understanding*, not, more narrowly, at a kind of *knowledge*. To specify the distinctive sort of understanding that is the characteristic goal of philosophers’ writing is to say what distinguishes that enterprise from that of other sorts of constructive seekers of understanding, such as novelists and scientific theorists. I want to do so by focusing not on the peculiar genre of nonfiction creative writing by which philosophical understanding is typically conveyed (though I think that subject is worthy of consideration), but rather on what is distinctive about the understanding itself: both its particular topic, and its characteristic goal. Philosophy is a self-reflexive enterprise: understanding is not only the *goal* of philosophical inquiry, but its *topic* as well. *We* are its topic; but it is us specifically as *understanding* creatures: *discursive* beings, makers and takers of *reasons*, seekers and speakers of *truth*. Seeing philosophy as addressing the nature and conditions of our rationality is, of course, a very traditional outlook—so traditional, indeed, that it is liable to seem quaint and oldfashioned. I’ll address this issue later, remarking now only that rationalism is one thing, and intellectualism another: pragmatists, too, are concerned with the practices of giving and asking for reasons.

²I am describing, of course, for the concept philosophy an exercise of the sort of recollective rationality (Hegel’s “Erinnerung”) considered for ordinary determinate empirical concepts in Chapter Three.

I understand the task of philosophers to have as a central element the explication of concepts—or, put slightly more carefully, the development and application of expressive tools with which to make explicit what is implicit in the use of concepts. When I say "explication of concepts", it is hard not to hear "analysis of meanings." There are obviously affinities between my specification and that which defined the concern specifically of "analytic philosophy" in the middle years of this century. Indeed, I intend, *inter alia*, to be saying what was right about that conception. But what I have in mind is different in various ways. *Explication*, making explicit, is not the same as *analysis*, at least as that notion was classically conceived. As I use the term, for instance, we have no more privileged access to the contents of our concepts than we do to the facts we use them to state; the concepts and the facts are two sides of one coin.

But the most important difference is that where analysis of meanings is a fundamentally *conservative* enterprise (consider the paradox of analysis), I see the point of explicating concepts rather to be opening them up to rational *criticism*. The rational enterprise, the practice of giving and asking for reasons that lies at the heart of discursive activity, requires not only criticizing *beliefs*, as false or unwarranted, but also criticizing *concepts*. Defective concepts distort our thought and constrain us by limiting the propositions and plans we can entertain as candidates for endorsement in belief and intention. This constraint operates behind our backs, out of our sight, since it limits what we are so much as capable of being aware of. Philosophy, in developing and applying tools for the rational criticism of concepts, seeks to free us from these fetters, by bringing the distorting influences out into the light of conscious day, exposing the commitments implicit in our concepts as vulnerable to rational challenge and debate.

2. The first thing to understand about concepts is that concept is a *normative* concept. This is a lesson we owe ultimately to Kant—the great, gray mother of us all. Kant saw us above all as traffickers in concepts. In fact, in a strict sense, *all* that kantian rational creatures can do is to apply concepts. For that is the genus he took to comprise both *judgment* and *action*, our theoretical activity and our practical activity. One of Kant's great innovations was his view that what in the first instance distinguishes judgments and actions from the mere behavior of denizens of the realm of nature is that they are things that we are in a distinctive sense *responsible* for. They express *commitments* of ours. The norms or rules that determine what we have committed ourselves to, what we have made ourselves responsible for, by making a judgment or performing an action, Kant calls 'concepts'. Judging and acting involves undertaking commitments whose credentials are always potentially at issue. That is, the commitments embodied in judgments and actions are ones we may or may not be *entitled* to, so that the question of whether they are *correct*, whether they are commitments we *ought* to acknowledge and embrace, can always be raised. One of the forms taken by the responsibility we undertake in judging and acting is the responsibility to give reasons that justify the judgment or the action. And the rules that are the concepts we apply in judging and acting determine what would count as a

reason for the judgment and the action.

Commitment, entitlement, responsibility—these are all normative notions. Kant replaces the *ontological* distinction between the physical and the mental with the *deontological* distinction between the realm of nature and the realm of freedom: the distinction between things that merely act regularly and things that are subject to distinctively normative sorts of assessment.

Thus for Kant the great philosophical questions are questions about the source and nature of normativity—of the bindingness or validity [Gültigkeit] of conceptual rules. Descartes had bequeathed to his successors a concern for *certainty*: a matter of our grip on concepts and ideas—paradigmatically, whether we have a hold on them that is clear and distinct. Kant bequeaths to his successors a concern rather for *necessity*: a matter of the grip concepts have on us, the way they bind or oblige us. ‘Necessary’ [notwendig] for Kant just means “according to a rule”. (That is why he is willing to speak of moral and natural necessity as species of a genus.) The important lesson he takes Hume to have taught isn’t about the threat of skepticism, but about how empirical knowledge is unintelligible if we insist on merely *describing* how things in fact *are*, without moving beyond that to *prescribing* how they *must* be, according to causal rules, and how empirical motivation (and so agency) is unintelligible if we stay at the level of ‘*is*’ and eschew reference to the ‘*ought*’s that outrun what merely is. Looking farther back, Kant finds “the celebrated Mr. Locke” sidetracked into a mere “physiology of the understanding”—the tracing of causal antecedents of thought in place of its justificatory antecedents—through a failure to appreciate the essentially normative character of claims to knowledge. But Kant takes the whole Enlightenment to be animated by an at least implicit appreciation of this point. For mankind’s coming into its intellectual and spiritual majority and maturity consists precisely in taking the sort of personal responsibility for its commitments, both doxastic and practical, insisted upon already by Descartes’ meditator.

This placing of normativity at the center of philosophical concern is the reason behind another of Kant’s signal innovations: the pride of place he accords to *judgment*. In a sharp break with tradition, he takes it that the smallest unit of experience, and hence of awareness, is the judgment. This is because judgments, applications of concepts, are the smallest unit for which knowers can be *responsible*. Concepts by themselves don’t express commitments; they only determine what commitments would be undertaken if they were applied. (Frege will express this kantian point by saying that judgeable contents are the smallest unit to which pragmatic force—paradigmatically the assertional force that consists in the assertor undertaking a special kind of commitment—can attach. Wittgenstein will distinguish sentences from terms and predicates as the smallest expressions whose free-standing utterance can be used to make a move in a language game.) The most general features of Kant’s understanding of the form of judgment also derive from its role as a unit of responsibility. The “I think” that can accompany all representations (hence being, in its formality, the emptiest

of all) is the formal shadow of the transcendental unity of apperception, the locus of responsibility determining a coresponsibility class of concept-applications (including actions), what is responsible *for* its judgments. The objective correlate of this subjective aspect of the form of judgment is the “object=X” to which the judgment is directed, the formal shadow of what the judgment makes the knower responsible *to*.

I think that philosophy is the study of us as creatures who judge and act, that is, as discursive, concept-using creatures. And I think that Kant is right to emphasize that understanding what we do in these terms is attributing to us various kinds of normative status, taking us to be subject to distinctive sorts of normative appraisal. So a central philosophical task is understanding this fundamental normative dimension within which we dwell. Kant’s own approach to this issue, developing themes from Rousseau, is based on the thought that genuinely normative authority (constraint by norms) is distinguished from causal power (constraint by facts) in that it binds only those who *acknowledge* it as binding. Because one is subject only to that authority one subjects oneself to, the normative realm can be understood equally as the realm of *freedom*. So being constrained by norms is not only compatible with freedom—properly understood, it can be seen to be what freedom consists in. I don’t know of a thought that is deeper, more difficult, or more important than this.

3. Kant’s most basic idea, I said, is that judgment and action are things we are in a distinctive way *responsible* for. What does it mean to be responsible for them? I think the kind of responsibility in question should be understood to be task responsibility: the responsibility to do something. What (else) do judging and acting oblige us to do? The commitments we undertake by applying concepts in particular circumstances—by judging and acting—are ones we may or may not be entitled to, according to the rules (norms) implicit in those concepts. Showing that we are entitled by the rules to apply the concept in a particular case is *justifying* the commitment we undertake thereby, offering *reasons* for it. That is what we are responsible for, the practical content of our conceptual commitments. In undertaking a conceptual commitment one renders oneself in principle liable to demands for reasons. The normative appraisal to which we subject ourselves in judging and acting is appraisal of our reasons. Further, offering a reason for the application of a concept is always applying another concept: making or rehearsing another judgment or undertaking or acknowledging another practical commitment (Kant’s “adopting a maxim”). Conceptual commitments both serve as and stand in need of reasons. The normative realm inhabited by creatures who can judge and act is not only the realm of freedom, it is the realm of reason.³

Understanding the norms for correct application that are implicit in concepts requires understanding the role those concepts play in reasoning: what (applications of concepts) count as reasons for the application of that concept, and what (applications of concepts) the application of that concept counts as a reason for. For apart from such understanding, one cannot fulfill the responsibility one

³This story is told in more detail in Chapter One.

undertakes by making a judgment or performing an action. So what distinguishes concept-using creatures from others is that we know our way around the *space of reasons*. Grasping or understanding a concept just is being able practically to place it in a network of inferential relations: to know what is evidence for or against its being properly applied to a particular case, and what its proper applicability to a particular case counts as evidence for or against. Our capacity to know (or believe) *that* something is the case depends on our having a certain kind of know *how*: the ability to tell what is a reason for what.

The cost of losing sight of this point is to assimilate genuinely conceptual activity, judging and acting, too closely to the behavior of mere animals—creatures who do not live and move and have their being in the normative realm of freedom and reason. We share with other animals (and for that matter, with bits of automatic machinery) the capacity reliably to respond differentially to various kinds of stimuli. We, like they, can be understood as classifying stimuli as being of certain kinds, insofar as we are disposed to produce different repeatable sorts of responses to those stimuli. We can respond differentially to red things by uttering the noise “That is red.” A parrot could be trained to do this, as pigeons are trained to peck at a different button when shown a red figure than when shown a green one. The empiricist tradition is right to emphasize that our capacity to have empirical knowledge begins with and crucially depends on such reliable differential responsive dispositions. But though the story begins with this sort of classification, it does not end there. For the rationalist tradition is right to emphasize that our classificatory responses count as applications of concepts, and hence as so much as candidates for knowledge, only in virtue of their role in reasoning. The crucial difference between the parrot’s utterance of the noise “That is red,” and the (let us suppose physically indistinguishable) utterance of a human reporter is that for the latter, but not the former, the utterance has the practical significance of making a claim. Doing that is taking up a normative stance of a kind that can serve as a premise from which to draw conclusions. That is, it can serve as a reason for taking up other stances. And further, it is a stance that itself can stand in need of reasons, at least if challenged by the adoption of other, incompatible stances. Where the parrot is merely responsively sounding off, the human counts as applying a concept just insofar as she is understood as making a move in a game of giving and asking for reasons.

The most basic point of Sellars’ rationalist critique of empiricism in his masterwork “Empiricism and the Philosophy of Mind,” is that even the *noninferentially* elicited perceptual judgments that the empiricist rightly appreciates as forming the empirical basis for our knowledge can count as judgments (applications of concepts) at all only insofar as they are *inferentially* articulated. Thus the idea that there could be an autonomous language game (a game one could play though one played no other) consisting entirely of noninferentially elicited reports—whether of environing stimuli or of the present contents of one’s own mind—is a radical mistake. To apply any concepts *noninferentially*, one must be able also to apply concepts inferentially. For it is an essential feature of concepts

that their applications can both serve as and stand in need of reasons. Making a report or a perceptual judgment is doing something that essentially, and not just accidentally, has the significance of making available a premise for reasoning. Learning to observe requires learning to infer. Experience and reasoning are two sides of one coin, two capacities presupposed by concept use that are in principle intelligible only in terms of their relations to each other.⁴

To claim that what distinguishes specifically conceptual classification from classification merely by differential responsive disposition is the inferential articulation of the response—that applications of concepts are essentially what can both serve as and stand in need of reasons—is to assign the game of giving and asking for reasons a preeminent place among discursive practices. For it is to say that what makes a practice *discursive* in the first place is that it incorporates reason-giving practices. Now of course there are many things one can do with concepts besides using them to argue and to justify. And it has seemed perverse to some post-Enlightenment thinkers in any way to privilege the rational, cognitive dimension of language use. But if the tradition I have been sketching is right, the capacity to use concepts in all the other ways explored and exploited by the artists and writers whose imaginative enterprises have rightly been admired by romantic opponents of logocentrism is parasitic on the prosaic inferential practices in virtue of which we are entitled to see concepts as in play in the first place. The game of giving and asking for reasons is not just one game among others one can play with language. It is the game in virtue of the playing of which what one has qualifies as *language* (or thought) at all. I am here disagreeing with Wittgenstein, when he claims that “language has no downtown.” On my view, it does, and that downtown (the region around which all the rest of discourse is arrayed as dependent suburbs, is the practices of giving and asking for reasons. This is a kind of linguistic *rationalism*. ‘Rationalism’ in this sense does not entail intellectualism, the doctrine that every *implicit* mastery of a propriety of practice is ultimately to be explained by appeal to a prior *explicit* grasp of a principle. It is entirely compatible with the sort of pragmatism that sees things the other way around.

4. As I am suggesting that we think of them, concepts are broadly inferential norms that implicitly govern practices of giving and asking for reasons. Dummett has suggested a useful model for thinking about the inferential articulation of conceptual contents. Generalizing from the model of meaning Gentzen introduces for sentential operators, Dummett suggests that we think of the use of any expression as involving two components: the circumstances in which it is appropriately used and the appropriate consequences of such use. Since our concern is with the application of the concepts expressed by using linguistic expressions, we can render this as the circumstances of appropriate application of the concept, and the appropriate consequences of such application—that is, what follows from the concept’s being applicable.

Some of the circumstances and consequences of applicability of a concept may

⁴Chapter Seven develops this theme further.

be inferential in nature. For instance, one of the circumstances of appropriate application of the concept red is that this concept is applicable wherever the concept scarlet is applicable. And to say that is just another way of saying that the inference from “X is scarlet,” to “X is red,” is a good one. And similarly, one of the consequences of the applicability of the concept red is the applicability of the concept colored. And to say that is just another way of saying that the inference from “X is red,” to “X is colored,” is a good one. But concepts like red also have *noninferential* circumstances of applicability, such as the visible presence of red things. And concepts such as unjust have noninferential consequences of application—that is, they can make it appropriate to *do* (or not do) something, to make another claim true, not just to *say* or judge that it is true.

Even the immediately empirical concepts of *observables*, which have noninferential *circumstances* of application and the immediately practical *evaluative* concepts, which have noninferential *consequences* of application, however, can be understood to have contents that are inferentially articulated. For all concepts incorporate an implicit commitment to the propriety of the inference from their circumstances to their consequences of application. One cannot use the concept red as including the circumstances and consequences mentioned above without committing oneself to the correctness of the inference from “X is scarlet,” to “X is colored.” So we might decompose the norms that govern the use of concepts into three components: circumstances of appropriate application, appropriate consequences of application, and the propriety of an inference from the circumstances to the consequences. I would prefer to understand the inferential commitment expansively, as including the circumstances and consequences it relates, and so as comprising all three normative elements.

I suggested at the outset that we think of philosophy as charged with producing and deploying tools for the criticism of concepts. The key point here is that concepts may incorporate defective inferences. Dummett offers this suggestive example:

A simple case would be that of a pejorative term, e.g. ‘Boche’. The conditions for applying the term to someone is that he is of German nationality; the consequences of its application are that he is barbarous and more prone to cruelty than other Europeans. We should envisage the connections in both directions as sufficiently tight as to be involved in the very meaning of the word: neither could be severed without altering its meaning. Someone who rejects the word does so because he does not want to permit a transition from the grounds for applying the term to the consequences of doing so.⁵

(It is useful to focus on a French epithet from the first world war, because we are sufficiently removed from its practical effect to be able to get a theoretical grip on how it works. But the thought should go over *mutatis mutandis* for

⁵Dummett, Frege: Philosophy of Language [Harper and Row, New York, 1973] p. 454.

pejoratives in current circulation.) Dummett's idea is that if you do not accept as correct the inference from German nationality to an unusual disposition to barbarity and cruelty, you can only reject the word. You cannot deny that there are any Boche, for that is just denying that the circumstances of application are ever satisfied, that is, that there are any Germans. And you cannot admit that there are Boche but deny that they are disposed to barbarity and cruelty (this is the "Some of my best friends are Boche," ploy), since that is just taking back in one breath what one has asserted just before. Any use of the term commits the user to the inference that is curled up, implicitly, in it. (At Oscar Wilde's trial the prosecutor read out some passages from the Importance of Being Earnest and said "I put it to you, Mr. Wilde, that this is blasphemy. Is it? Yes or no?" Wilde replied just as he ought on the account I am urging: "Sir, 'blasphemy' is not one of my words."⁶)

Although they are perhaps among the most dangerous, it is not just highly-charged words, words that couple 'descriptive' circumstances of application with 'evaluative' consequences of application that incorporate inferences of which we may need to be critical. The use of *any* expression involves commitment to the propriety of the inference from its circumstances to its consequences of application. These are almost never logically valid inferences. On the contrary, they are what Sellars called "material" inferences: inferences that articulate the content of the concept expressed. Classical disputes about the nature of personal identity, for instance, can be understood as taking the form of arguments about the propriety of such a material inference. We can agree, we may suppose, about the more or less forensic consequences of application of the concept "same person," having in mind its significance for attributions of (co-)responsibility. When we disagree about the circumstances of application that should be paired with it—for instance whether bodily or neural continuity, or the psychological continuity of memory count for more—we are really disagreeing about the correctness of the inference from the obtaining of these conditions to the ascription of responsibility. The question about what is the correct concept is a question about which inferences to endorse. I think it is helpful to think about a great number of the questions we ask about other important concepts in these same terms: as having the form of queries about what inferences from circumstances to consequences of application we ought to acknowledge as correct, and why. Think in these terms about such very abstract concepts as morally wrong, just, beautiful, true, explain, know, or prove, and again about 'thicker' ones such as unkind, cruel, elegant, justify, and understand.

The use of any of these concepts involves a material inferential commitment: commitment to the propriety of a substantial inferential move from the circumstances in which it is appropriate to apply the concept to the consequences of doing so. The concepts are substantive just because the inferences they incorporate are. Exactly this commitment becomes invisible, however, if one conceives conceptual

⁶Of course, being right on this point didn't keep Wilde out of trouble, anymore than it did Salman Rushdie.

content in terms of *truth conditions*. For the idea of truth conditions is the idea of a single set of conditions that are at once necessary and sufficient for the application of the concept. The idea of individually necessary conditions that are also jointly sufficient is the idea of a set of consequences of application that can also serve as circumstances of application. Thus the circumstances of application are understood as already including the consequences of application, so that no endorsement of a substantive inference is involved in using the concept. The concept of concepts like this is not incoherent. It is the ideal of *logical* or *formal* concepts. Thus it is a criterion of adequacy for introducing logical connectives that they be inferentially conservative: that their introduction and elimination rules be so related that they permit no new inferences involving only the old vocabulary. But it is a bad idea to take this model of the relation between circumstances and consequences of application of logical vocabulary and extend it to encompass also the substantively contentful *nonlogical* concepts that are the currency in which most of our cognitive and practical transactions are conducted.

It is a bad idea because of its built-in conservatism. Understanding meaning or conceptual content in terms of truth conditions—individually necessary and jointly sufficient conditions—squeezes out of the picture the substantive inferential commitment implicit in the use of any nonlogical concept. But it is precisely those inferential commitments that are subject to *criticism* in the light of substantive collateral beliefs. If one does not believe that Germans are distinctively barbarous or prone to cruelty, then one must not use the concept Boche, just *because* one does not endorse the substantive material inference it incorporates. On the other model, this diagnosis is not available. The most one can say is that one does not know how to specify truth conditions for the concept. But just what is objectionable about it and why does not appear from this theoretical perspective. . Criticism of concepts is always criticism of the inferential connections. For criticizing whether all the individually sufficient conditions (circumstances) “go together”, i.e. are circumstances of application of one concept, just is wondering whether they all have the same consequences of application (and similarly for wondering whether the consequences of application all “go together”).

5. When we think of conceptual contents in the way I am recommending, we can see not only how beliefs can be used to criticize concepts, but also how concepts can be used to criticize beliefs. For it is the material inferences incorporated in our concepts that we use to elaborate the antecedents and consequences of various candidates for belief—to tell what we would be committing ourselves to, what would entitle us to those commitments, what would be incompatible with them, and so on. Once it is accepted that the inferential norms implicit in our concepts are in principle as revisable in the light of evidence as particular beliefs, conceptual and empirical authority appear as two sides of one coin. Rationally justifying our concepts depends on finding out about how things are—about what actually follows from what—as is most evident in the case of massively defective concepts such as Boche.

Adjusting our beliefs in the light of the connections among them dictated by our concepts, and our concepts in the light of our evidence for the substantive beliefs presupposed by the inferences they incorporate, is the rationally reflective enterprise introduced to us by Socrates. It is what results when the rational, normative connections among claims that govern the practice of giving and asking for reasons are themselves brought into the game, as liable to demands for reasons and justification. Saying or thinking something, making it explicit, consists in applying concepts, thereby taking up a stance in the space of reasons, making a move in the game of giving and asking for reasons. The structure of that space, of that game, though, is not given in advance of our finding out how things are with what we are talking about. For what is *really* a reason for what depends on how things *actually* are. But that inferential structure itself can be the subject of claims and thoughts. It can itself be made explicit in the form of claims about what follows from what, what claims are evidence for or against what other claims, what else one would be committing oneself to by making a certain judgment or performing a certain action. So long as the commitment to the propriety of the inference from German nationality to barbarity and unusual cruelty remains merely implicit in the use of term such as ‘Boche’, it is hidden from rational scrutiny. When it is made explicit in the form of the conditional claim “Anyone who is German is barbarous and unusually prone to cruelty,” it is subject to rational challenge and assessment; it can, for instance, be confronted with such counterexamples as Bach and Goethe.

Discursive explicitness, the application of concepts, is Kantian apperception or consciousness. Bringing into discursive explicitness the inferentially articulated conceptual norms in virtue of which we can be conscious or discursively aware of anything at all is the task of reflection, or self-consciousness. This is the expressive task distinctive of philosophy. Of course, the practitioners of special disciplines, such as membrane physiology, are concerned to unpack and criticize the inferential commitments implicit in using concepts such as lipid soluble with a given set of circumstances and consequences of application, too. It is the emphasis on the “anything at all” distinguishes philosophical reflection from the more focused reflection that goes on within such special disciplines. Earlier I pinned on Kant a view that identifies us as distinctively *rational* creatures, where that is understood as a matter of our being subject to a certain kind of *normative* assessment: we are creatures who can undertake *commitments* and *responsibilities* that are *conceptually* articulated in that their contents are articulated by what would count as *reasons* for them (as well as what other commitments and responsibilities they provide reasons for). One of philosophy’s defining obligations is to supply and deploy an expressive toolbox, filled with concepts that help us make explicit various aspects of *rationality* and *normativity* in general. **The topic of philosophy is normativity in all its guises, and inference in all its forms.** And its task is an *expressive, explicative* one. So it is the job of practitioners of the various philosophical subfields to design and produce specialized expressive tools, and to hone and shape them with use. At the most general level, *inferential* connections are made explicit by *conditionals*, and

their *normative* force is made explicit by *deontic* vocabulary. Different branches of philosophy can be distinguished by the different sorts of inference and normativity they address and explicate, the various special senses of “if... then_____,” or of ‘ought’ for which they care. Thus philosophers of science, for instance, develop and deploy conditionals codifying causal, functional, teleological, and other explanatory inferential relations, value theorists sharpen our appreciation of the significance of the differences in the endorsements expressed by prudential, legal, ethical, and aesthetic ‘ought’s, and so on.

6. I said at the outset that I thought of philosophy as defined by its history, rather than by its nature, but that, following Hegel, I think of our task as understanding it by finding or making a nature in or from its history. The gesture I have made in that direction today, though, could also be summarized in a different kind of definition, namely in the ostensive definition: Philosophy is the kind of thing that Kant and Hegel did (one might immediately want to add Plato, Aristotle, Frege and Wittgenstein to the list, and then we are embarked on the enterprise of turning a gesture into a story, indeed, a history). So one might ask: Why not just say that, and be done with it? While, as I’ve indicated, I think that specification is a fine place to start, I also think there is a point to trying to be somewhat more explicit about just what sort of thing it is that one takes it Kant and Hegel (and Frege and Wittgenstein) did. Doing that is not being satisfied just with a wave at philosophy as something that has a history. It is trying rationally to reconstruct that tradition, to recast it into a form in which a constellation of ideas can be seen to be emerging, being expressed, refined, and developed.

With those giants, I see philosophy as a discipline whose distinctive concern is with a certain kind of *self-consciousness*: awareness of ourselves as specifically *discursive* (that is, concept-mongering) creatures. Its task is understanding the conditions, nature, and consequences of conceptual norms and the activities—starting with the social practices of giving and asking for reasons—that they make possible and that make them possible. As concept users, we are beings who can make explicit how things are and what we are doing—even if always only in relief against a background of implicit circumstances, conditions, skills, and practices. Among the things on which we can bring our explicating capacities to bear are those very concept-using capacities that make it possible to make anything at all explicit. Doing that, I am saying, is philosophizing.

It is easy to be misled by the homey familiarity of these sentiments, and correspondingly important to distinguish this characterization from some neighbors with which it is liable to be confused. There is a clear affinity between this view and Kant’s coronation of philosophy as “queen of the sciences.” For on this account philosophy does extend its view to encompass all activity that is discursive in a broad sense—that is, all activity that presupposes a capacity for judgment and agency, sapience in general. But in this sense, philosophy is at most *a* queen of the sciences, not *the* queen. For the magisterial sweep of its purview does not serve to distinguish it from, say, psychology, sociology, history, literary or

cultural criticism, or even journalism. What distinguishes it is the *expressive* nature of its concern with discursiveness in general, rather than its inclusive scope. My sketch was aimed at introducing a specific difference pertaining to philosophy, not a unique privilege with respect to such other disciplines.

Again, as I have characterized it, philosophy does not play a *foundational* role with respect to other disciplines. Its claims do not stand prior to those of the special sciences in some order of ultimate justification. Nor does philosophy sit at the other end of the process as final judge over the propriety of judgments and actions—as though the warrant of ordinary theoretical and practical applications of concepts remained somehow provisional until certified by philosophical investigation. And philosophy as I have described it likewise asserts no methodological privilege or insight that potentially collides with the actual procedures of other disciplines.

Indeed, philosophy's own proper concerns with the nature of normativity in general, and with its conceptual species in particular, so on inference and justification in general, impinge on the other disciplines in a role that equally well deserves the characterization of "handmaiden." For what we do that has been misunderstood as having foundational or methodological significance is provide and apply tools for unpacking the substantive commitments that are implicit in the concepts deployed throughout the culture, including the specialized disciplines of the high culture. Making those norms and inferences explicit in the form of claims exposes them for the first time to reasoned assessment, challenge, and defense, and so to the sort of rational emendation that is the primary process of conceptual evolution. But once the implicit presuppositions and consequences have been brought out into the daylight of explicitness, the process of assessment, emendation, and so evolution is the business of those whose concepts they are—and not something philosophers have any particular authority over or expertise regarding. Put another way, it is the business of philosophers to figure out ways to increase semantic and discursive self-consciousness. What one does with that self-consciousness is not our business *qua* philosophers—though of course, *qua* intellectuals generally, it may well be.

Philosophy's *expressive* enterprise is grounded in its focus on us as a certain kind of thing, an expressing thing: as at once creatures and creators of conceptual norms, producers and consumers of reasons, beings distinguished by being subject to the peculiar normative force of the better reason. Its concern with us as specifically *normative* creatures sets philosophy off from the empirical disciplines, both the natural and the social sciences. It is this normative character that binds together the currents of thought epitomized in Stanley Cavell's characteristically trenchant aphorism that Kant depsychologized epistemology, Frege depsychologized logic, and Wittgenstein depsychologized psychology. We might add that Hegel depsychologized history. The depsychologizing move in question is equally a desociologizing. For it is a refocusing on the *normative bindingness* of the concepts deployed in ground-level empirical knowledge, reasoning, and thought in general. This is a move beyond the narrowly *natural* (in the sense of

the describable order of causes), toward what Hegel called the ‘spiritual’ [geistig], that is, the *normative* order. That its concern is specifically with our *conceptual* normativity sets philosophy off from the other humanistic disciplines, from the literary as well as the plastic arts. Conceptual commitments are distinguished by their inferential articulation, by the way they can serve as reasons for one another, and by the way they stand in need of reasons, their entitlement always potentially being at issue. Now in asserting the centrality and indispensability, indeed, the criterial role, of practices of giving and asking for reasons, I am far from saying that reasoning—or even thinking—is all anyone ought to do. I am saying that philosophers’ distinctive concern is with what else those reason-mongering practices make possible, and how they do, on the one hand, and with what it is that makes them possible—what sort of doings count as sayings, how believing or saying that is founded on knowing how—on the other. It is this distinctive constellation of concerns that makes philosophy the party of reasons, and philosophers the friends of the norms, the ones who bring out into the light of discursive explicitness our capacity to make things discursively explicit.